

CSE201: Monsoon 2024

Advanced Programming

Lecture 17: Mutual Exclusion

Dr. Arun Balaji Buduru

Founding Head, Usable Security Group (USG)


Associate Professor, Dept. of CSE | HCD

IIIT-Delhi, India

Today's Lecture

- Race conditions
- Mutual exclusion
- Monitor locks
- Memory consistency
- Producer consumer problem

Race Condition



Put green pieces

How can we have
alternating colors?

Put red pieces


Mutual Exclusion

- ***Critical section:*** a block of code that access shared modifiable data or resource that should be operated on by only one thread at a time
- ***Mutual exclusion:*** a property that ensures that a critical section is only executed by a thread at a time.
 - *Otherwise it results in a race condition!*



Threads

- A *thread* A is (formally) a sequence a_0, a_1, \dots of events
 - Notation: $a_0 \rightarrow a_1$ indicates order



Example Thread Events

- Assign to shared variable
- Assign to local variable
- Invoke method
- Return from method
- Lots of other things ...

Concurrent Execution Over Multiple Threads

- Thread A




- Thread B



Interleavings

- Events of two or more threads
 - Interleaved
 - Not necessarily independent (why?)



Question

```
class Counter implements Runnable {
    int counter = 0;

    public void run() { counter++; }
    public static void main(String[] args)
        throws InterruptedException {
        ExecutorService exec =
            Executors.newFixedThreadPool(2);

        Counter task = new Counter();
        for(int i=0; i<1000; i++) {
            exec.execute(task);
        }

        if(!exec.isTerminated()) {
            exec.shutdown();
            exec.awaitTermination(5L, TimeUnit.SECONDS);
        }

        System.out.println(task.counter);
    }
}
```

- What will be the output of this program?
 - Race on counter!
 - Buggy code and you will see different answers in different runs

Implementing Mutual Exclusion

```
class Counter implements Runnable {
    int counter = 0;
    // Both the versions of run method below is correct
    public synchronized void run() { counter++; }
    /* public void run() { synchronized(this) {counter++;} } */
    public static void main(String[] args)
        throws InterruptedException {
        ExecutorService exec =
            Executors.newFixedThreadPool(2);

        Counter task = new Counter();
        for(int i=0; i<1000; i++) {
            exec.execute(task);
        }


        if(!exec.isTerminated()) {
            exec.shutdown();
            exec.awaitTermination(5L, TimeUnit.SECONDS);
        }

        System.out.println(task.counter);
    }
}
```

- Critical section


- The synchronized methods (or block) define the critical sections
- By using synchronized keyword we achieved mutual exclusion
 - Now let's analyze this

Monitors



- Each object has a “**monitor**” that is a token used to determine which application **thread** has control of a particular **object** instance
- In execution of a synchronized method (or block), access to the object monitor must be gained before the execution
- Access to the object monitor is queued
- Entering a monitor is also referred to as **locking** the monitor, or **acquiring ownership** of the monitor
- If a thread *A* tries to acquire ownership of a monitor and a different thread has already entered the monitor, the current thread (*A*) must wait until the other thread leaves the monitor

Analyzing our Counter Increment Example



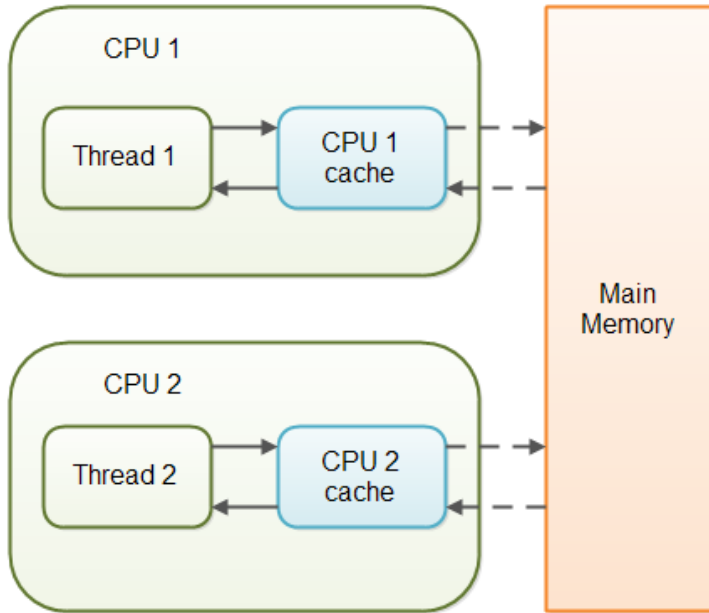
- Only one thread can get the “key” to enter the “run” method i.e., take a lock on monitor
- Rest all threads will be queued to get the lock on monitor
- **Note:** There is no guarantee for fairness, i.e. longest waiting thread need not always get the lock first

We are Still Missing Something...

```
class Counter implements Runnable {
    static int counter = 0;
    static int turn = RED; //finals RED=0 and GREEN=1
    int me, other;
    public Counter(int c1, int c2) { me=c1; other=c2; }
    synchronized static void update(int me, int other) {
        if(counter<MAX && turn==me) {
            counter++; turn=other;
        }
    }
    public void run() {
        while(counter < MAX) {
            if(turn == me) {
                update(me, other);
            }
        }
    }
    public static void main(String args[])throws InterruptedException{
        Counter task1 = new Counter(RED, GREEN);
        Counter task2 = new Counter(GREEN, RED);
        Thread t1 = new Thread(task1); Thread t2 = new Thread(task2);
        t1.start(); t2.start(); t1.join(); t2.join();
    }
}
```

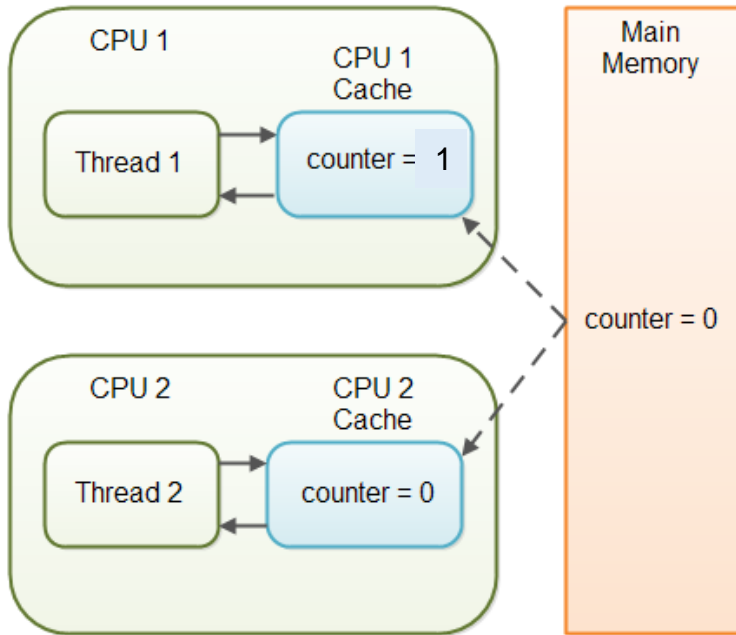
- This program will never terminate
- Using synchronized is just one part of the perfect solution
- Although there is no race on shared variables **counter** and **color**, the value of **counter** and **color** that a thread begins with may not be its last updated value

Memory Consistency Issue (1/2)



- Modern computing systems use multicore processors
- Each core has its own local cache
- For faster data access, memory referenced by a CPU is first copied from main memory (RAM) onto its local cache
- The updated memory content on cache is not immediately written back to RAM
 - This memory address might be referenced again in near future, hence immediately writing the cache content to RAM can hamper performance

Memory Consistency Issue (2/2)



- Imagine Counter example has two threads in its thread pool – Thread 1 on CPU1 and Thread 2 on CPU2
- Thread 1 increments counter from 0 to 1. This updated value resides on the cache of CPU1 and might not be immediately written back to the RAM
- Thread 2 now gets the chance to update the counter. It fetches the counter content from RAM but this is the old value (=0) and not the last updated value (=1)
- This is memory consistency error!

The Correct Version of Counter Code

```
class Counter implements Runnable {  
    volatile static int counter = 0;  
    volatile static int turn = RED;  
  
    .....  
    .....  
}
```

- Declare the counter as “**volatile**”
- Indication to JVM for storing the value of counter & color on RAM after every update to it
- With this each thread will always get the latest value of the counter & color

Creating an Object Lock

```
class Counter implements Runnable {  
    volatile int counter = 0;  
  
    private Object lock = new Object();  
    public void run() {  
        synchronized(lock) {  
            counter++;  
        }  
    }  
  
    .....  
    .....  
}
```

- We can also pass any object instance to synchronized

Monitor Locks are Reentrant

```
class Counter implements Runnable {  
    volatile int counter = 0;  
  
    public synchronized int value() { return counter; }  
  
    public synchronized void run() {  
        if(value() < 100) {  
            counter++;  
        }  
    }  
  
    .....  
    .....  
}
```

- Both value() and run() are synchronized methods
- Monitor locks are **reentrant** in Java
 - Same thread can recursively take the same lock
- Once a thread has taken a monitor lock then any further request by this same thread to reacquire the same monitor lock is redundant
- Monitor lock is released only after exiting the oldest synchronized block

Demerits of Monitor Lock

- Does not guarantee fairness
 - Lock might not be given to the longest waiting thread
- Might lead to starvation
 - A thread can indefinitely hold the monitor lock for doing some big computation while other threads keep waiting to get this monitor lock
 - Not possible to interrupt the thread who owns the lock
 - Not possible for a thread to decline waiting for the lock if its unavailable

The Producer Consumer Problem

- We need to synchronize between transactions, for example, the consumer-producer scenario



Wait and Notify

- Allows two threads to cooperate
- Based on a single shared lock object
 - Marge put a cookie wait and notify Homer
 - Homer eat a cookie wait and notify Marge
 - Marge put a cookie wait and notify Homer
 - Homer eat a cookie wait and notify Marge

The `wait()` Method

- The `wait()` method is part of the class `java.lang.Object`
- It requires a lock on the object's monitor to execute
- It must be called from a synchronized method, or from a synchronized segment of code
- `wait()` causes the current thread to relinquish the CPU and wait until another thread invokes the `notify()` method or the `notifyAll()` method for this object
- Upon call for `wait()`, the thread releases ownership of this monitor and waits until another thread notifies the waiting threads of the object

Wait/Notify Sequence

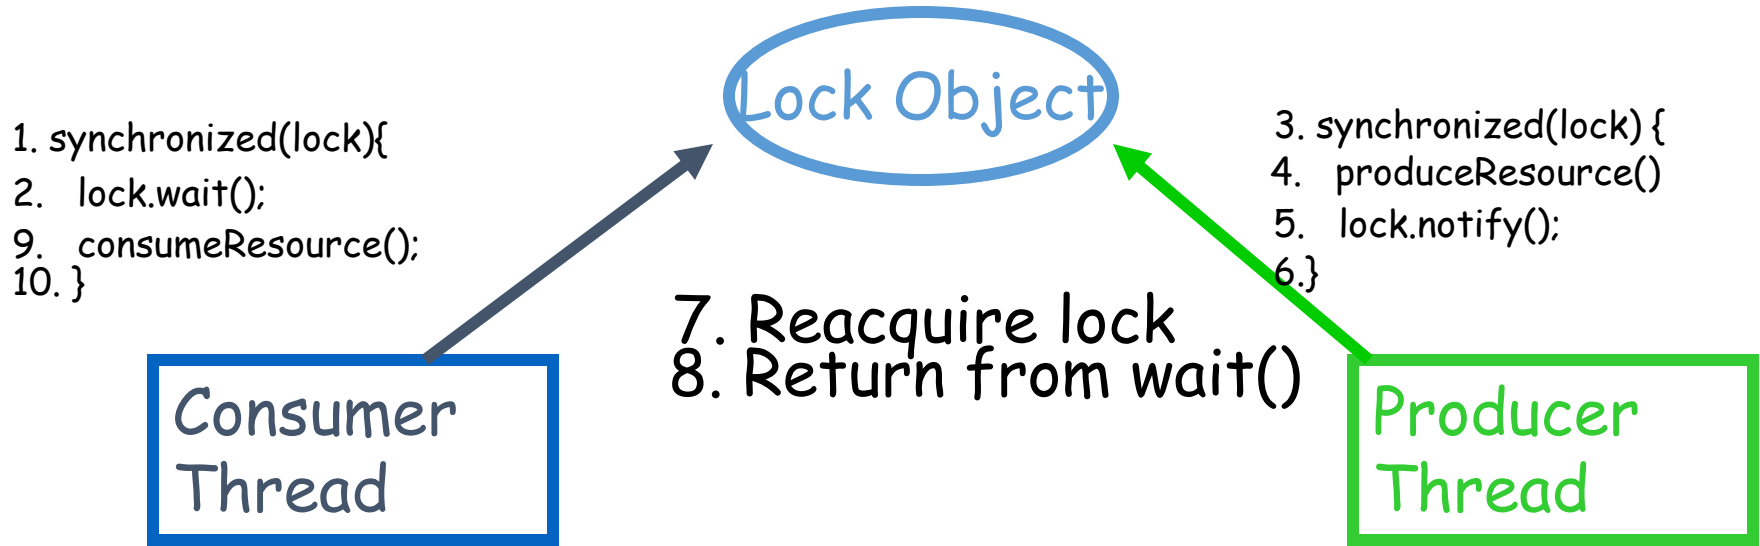
```
1. synchronized(lock){  
2.  lock.wait();  
9.  consumeResource();  
10. }
```

Consumer


```
3. synchronized(lock) {  
4.  produceResource()  
5.  lock.notify();  
6. }
```

Producer

Wait/Notify Sequence



Wait/Notify Sequence



Wait/Notify Sequence

```
1. synchronized(lock){  
2.  lock.wait();  
9.  consumeResource();  
10. }
```

Consumer
Thread

Lock Object

```
7. Reacquire lock  
8. Return from wait()
```


```
3. synchronized(lock) {  
4.  produceResource()  
5.  lock.notify();  
6. }
```

Producer
Thread

Wait/Notify Sequence

```
1. synchronized(lock){  
2.  lock.wait();  
9.  consumeResource();  
10. }
```

Consumer
Thread



```
7. Reacquire lock  
8. Return from wait()
```


```
3. synchronized(lock) {  
4.  produceResource()  
5.  lock.notify();  
6. }
```

Producer
Thread

Wait/Notify Sequence

```
1. synchronized(lock){  
2.  lock.wait();  
9.  consumeResource();  
10. }
```

Consumer
Thread

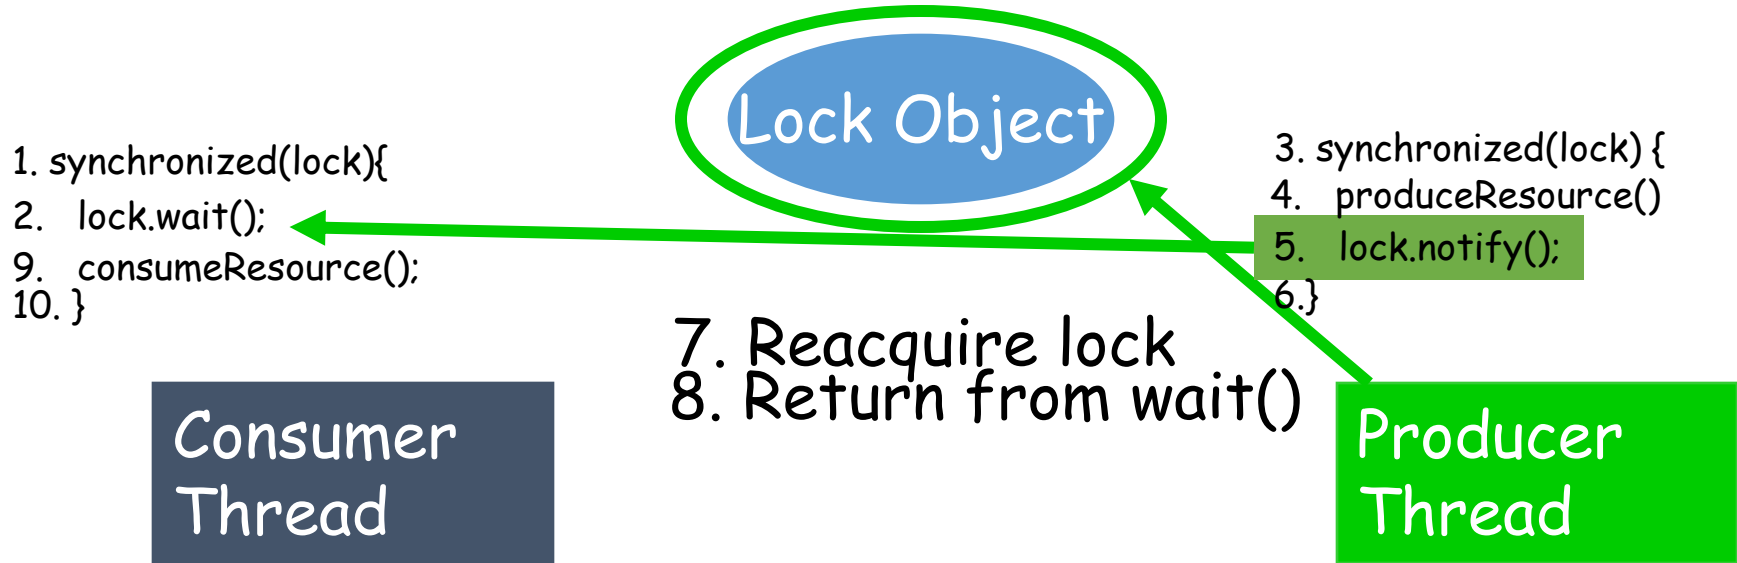


```
7. Reacquire lock  
8. Return from wait()
```

```
3. synchronized(lock) {  
4.  produceResource()  
5.  lock.notify();  
6. }
```

Producer
Thread

Wait/Notify Sequence



Wait/Notify Sequence

```
1. synchronized(lock){  
2.  lock.wait();  
9.  consumeResource();  
10. }
```


Consumer
Thread

Lock Object

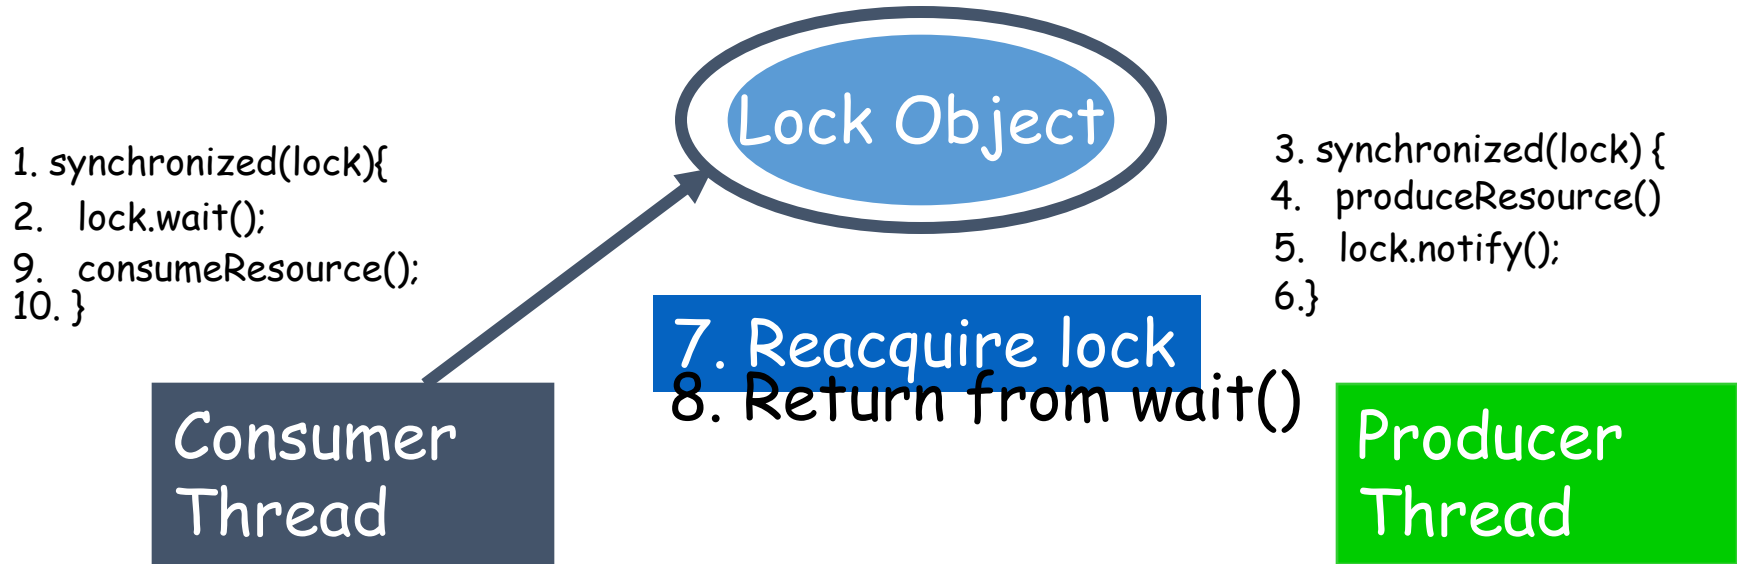
7. Reacquire lock
8. Return from wait()

```
3. synchronized(lock) {  
4.  produceResource()  
5.  lock.notify();  
6. }
```


Producer
Thread




Wait/Notify Sequence



Wait/Notify Sequence



Wait/Notify Sequence



Wait/Notify Sequence

```
1. synchronized(lock){  
2.  lock.wait();  
9.  consumeResource();  
10. }
```

Consumer
Thread

Lock Object

```
7. Reacquire lock  
8. Return from wait()
```

```
3. synchronized(lock) {  
4.  produceResource()  
5.  lock.notify();  
6. }
```

Producer
Thread

The Simpsons: Main Method

```
public class SimpsonsTest {  
    public static void main(String[] args) {  
        CookieJar jar = new CookieJar();  
        Homer homer = new Homer(jar);  
        Marge marge = new Marge(jar);  
        new Thread(homer).start();  
        new Thread(marge).start();  
    }  
}
```

The Simpsons: Homer

```
class Homer implements Runnable {  
    CookieJar jar;  
  
    public Homer(CookieJar jar) {  
        this.jar = jar;  
    }  
  
    public void eat() {  
        jar.getCookie("Homer");  
        try {  
            Thread.sleep((int)Math.random() * 500);  
        } catch (InterruptedException ie) {}  
    }  
  
    public void run() {  
        for (int i = 0 ; i < 5 ; i++) eat();  
    }  
}
```

The Simpsons: Marge

```
class Marge implements Runnable {  
    CookieJar jar;  
  
    public Marge(CookieJar jar) {  
        this.jar = jar;  
    }  
  
    public void bake(int cookieNumber) {  
        jar.putCookie("Marge", cookieNumber);  
        try {  
            Thread.sleep((int)Math.random() * 500);  
        } catch (InterruptedException ie) {}  
    }  
  
    public void run() {  
        for (int i = 0 ; i < 5 ; i++) bake(i);  
    }  
}
```

The Simpsons: CookieJar

```
class CookieJar {  
    private volatile int contents;  
    private volatile boolean available = false;  
  
    public synchronized void getCookie(String who)  
    {  
        while (!available) {  
            try {  
                wait();  
            } catch (InterruptedException e) { }  
        }  
        available = false;  
        notifyAll();  
        System.out.println( who + " ate cookie "  
                               +  
                               contents);  
    }  
}
```

```
    public synchronized void putCookie(String who,  
                                         int value) {  
        while (available) {  
            try {  
                wait();  
            } catch (InterruptedException e) { }  
        }  
        contents = value;  
        available = true;  
        System.out.println(who + " put cookie " +  
                           contents + " in the jar");  
        notifyAll();  
    }  
} /* end of class CookieJar */
```

The Simpsons: Output

Marge put cookie 0 in the jar

Homer ate cookie 0

Marge put cookie 1 in the jar

Homer ate cookie 1

Marge put cookie 2 in the jar

Homer ate cookie 2

Marge put cookie 3 in the jar

Homer ate cookie 3

Marge put cookie 4 in the jar

Homer ate cookie 4